

64.

Assumption : umask applies (we don't change it,

```
int make_new_log (char *fname) {
```

```
    struct stat sb;
```

```
    int rcs = -1;
```

Or access

or open.

```
    if (stat (fname, &sb) == -1) {
```

```
        rcs = open (fname, O_WRONLY | O_CREAT | O_TRUNC,
```

```
                (S_IRUSR | S_IWUSR | S_ISGRP |  
                 S_IWOTH));
```

}

```
    return rcs;
```

}

exkm not error

65.

```
# define SIZE 128
```

```
char *whereami () {
```

```
    char *buf = (char *) malloc (SIZE & sizeof (char)),
```

```
    int len = SIZE;
```

```
    getcmd (buf, len);
```

```
    while (!buf && errno == ERANGE) {
```

```
        len += SIZE;
```

```
        buf = (char *) realloc (buf, len);
```

```
        getcmd (buf, len);
```

}

/\* now clip & \*/

if (errno == 0 || errno == ERANGE)

buf = (char\*) realloc(buf, strlen(buf) + 1);

return buf;

}

66.

int isroot (const char \*path) {

struct stat sroot, sb;

if (stat("//&sroot") == -1) return 0;

if (stat(path, &sb) == -1) return 0;

if ((sroot.st\_ino == sb.st\_ino) &&

(sroot.st\_dev == sb.st\_dev)) return 1;

else return 0;

}

67.

int is\_leaf\_directory (const char \*path) {

struct stat sb;

DIR \*current;

struct dirent \*dir;

char pnd[pnd];

current = opendir(path);

if (!current) return 0;

pwd = getcwd(pwd, PATH\_MAX);

if (!pwd) return 0;

if (chdir(path) == -1) return 0;

while ((dirb = readdir(current)) != NULL)) {

if (strcmp(dirb->d\_name, ".") ||

strcmp(dirb->d\_name, "..")) continue;

if (stat(dirb->d\_name, &sb) == -1) return 0;

if (S\_ISDIR(sb.st\_mode)) return 0;

}

if (chdir(pwd) == -1) return 0;

return 1;

3

68. int rm\_tree(const char \*victim) {  
 Struct stat sb; if (stat(victim, &sb) == -1)  
 char pnd[PATH\_MAX]; return 0;  
 getcwd(pnd, PATH\_MAX); if (S\_ISDIR(sb.st\_mode))  
 if (!pnd) return 0;  
 DIR \*current;  
 Struct dirent dirb;  
 current = opendir(victim);  
 if (!current) return 0;  
 if (chkdir(victim) == -1) return 0;  
 while ((dirb = readdir(current)) != NULL) {  
 if (!strcmp(dirb->d\_name, ".") ||  
 !strcmp(dirb->d\_name, "..")) continue;  
 if (!rm\_tree(dirb->d\_name)) return 0;  
 }

if (closedir(current) == -1) return 0;  
 if (rmdir(victim) == -1) return 0;

else {

if (unlink(victim) == -1) return 0;

}

return 1;

3

```

int countfiles ( const charpath)
int res=0 struct stat sb; if (stat (victim, &sb) == -1)
    char pnd [PATH_MAX]; return 0;

getpid (pnd, PATH_MAX); if (S_ISDIR (sb.st_mode)) {
if (!pnd) return -1; int buf=0;
DIR *current;
struct dirent dirb;
current = opendir (victim);
if (!current) return -1;
if (chkdir (victim) == -1) return -1;
while ((dirb = readdir (current)) != NULL) {
if (!strcmp (dirb->d_name, ".") || !strcmp (dirb->d_name, "..")) continue;
if (buf == countfiles (dirb->d_name) == -1)
    result += buf; return -1;
if (chkdir (pnd) == -1) return -1;
if (closedir (current) == -1) return 0;
}

else {
    result += 1;
}
return result;

```

1.

Struct stat f1, f2;

```
if(stat(this, f1) == -1) return 0;  
if(stat(that, f2) == -1) return 0;  
if(f1.st-dev == f2.st-dev &  
f1.st-ino == f2.st-ino) return 1;  
else return 0;
```

2.

```
int conditide(const char *username){  
    struct passwd *pw = getpwnam(username+1);  
    if(!pw) return -1;  
    return chdir(pw->pw_dir);
```

3.

Program has 42774758 bytes

3.)

P A → moe → with size 1 → 42774758 iterations  
which include ~11 system calls  
(a lot of time)  
P B → larry → with size 8792 → 42774758 / 8792  
iterations and system calls (may less)

P C → curly → with getchar (buffered I/O) it  
needs less system calls (buffered  
input in the background) but still  
has 42774758 iterations to get  
each char from flat buffer.

(4.) No password is generated with one way algorithm  
↳ Sys admin can only look up hashed password  
or change the password.

(5.) Fundamental difference  $\rightarrow$  system calls can not be replaced (at least it's really hard). Library functions can be replaced. Because system calls directly invoke requests in the kernel. Lib functions are mostly abstractions of common programming tasks.

(6.) No it is not. Only root can change file uid of a process.  $\rightarrow$  File is created with uid/gid of process that creates it.

(7.) a) hard link links name to i-node number (i-nodes are only unique within a filesystem). Different devices can have same i-node - it symbolic however links a name to a name ( $\rightarrow$  works across filesystems because they are unique).

b) "hard links keep the universe together." e.g. the contents of the i-nodes stay the same (even if new link to i-node and old one removed). With a symbolic file can quickly be chaos. (files can be renamed and new files with that name created)  $\rightarrow$  symbolic points to wrong file

(8.) mv on a filesystem just needs to link new location to i-node and unlink old location. (contents of file are not moved). To mv across filesystems we need to move the contents of the file (because the i-node is only unique within a filesystem).

(9.) hard links Link a name to an i-node. i-nodes however are only unique within a filesystem. (different filesystems can have the same i-node). → can not link across filesystems.

(10.) For mv within filesystem we only need to link new location to i-node and unlink old one. (Data of file does not need to be moved). To different filesystem data needs to be moved. If we do not have read permission on file → we can't move it,

(11.) The ctime is the last time the i-node was changed. changing the m- & a-time which are stored in the i-node require means changing the i-node → c-time gets updated to when the change happened.

12. It is not! To open/create a file you need write & execute permission on the directory.  
→ same permissions are needed to delete a file.

→ you can delete.

same as 12.

14. System calls invoke direct requests to the kernel. recreating that behaviour is very hard / impossible. Lib functions are used to simplify common programming tasks  
→ can be easily replaced.

15. see 5.

16. can create a loop in the filesystem.

17. add a line on the mail file.

↳ access time → turned since

↳ modification time → new mail received

18. no more inodes available

to create a file a filesystem needs enough space & i-node index is limited - number of

- (13.) The "fastfile" has holes in it.
- (20.) Shared is a hashed version of the password that was generated via a one way algorithm (need pw & salt to get that).
- (21.) program is a bundle of instructions designed to perform certain tasks.  
process is a running instance of a program  
(that's where it gets interesting)
- (22.)  $m \rightarrow$  i-node link number  $\rightarrow$  file is i-node change time
- (23.) to m file we only need write/execute perm on dir (user certainly has that)

5

3

4

umask:      | 0 |    0 1 |    1 0 0

~umask:      0 1 0      | 0 0    0 1 |

&amp;

permission :      | | |      | 0 1 |    0 0 1  
mask

↳ permissions:      0 1 0      | 0 0    0 0 1

=>      - n -      r - -      - - x  
user      group      other

a)

b) open can only set permission of newly created file

↳ existing file keeps old permission.

35.

euid = 4     $\longleftrightarrow$     fuid = 4

egid = 7

uid match → only user permissions apply

$\Rightarrow$  no x perm for user  $\Rightarrow$  x fails

(perm denied)

euid = 4     $\nabla$     fuid = 12

egid = 7     $\longleftrightarrow$     fgid = 7

$\rightarrow$  uid no match  $\Rightarrow$  gid match  $\Rightarrow$  only group

permissions are looked at.  $\Rightarrow$  group has x perm

$\Rightarrow$  perm granted  $\Rightarrow$  x succeeds

36.

(37)

euid = 10      u      fuid = 12

egid = 7      u      fgid = 16

$\Rightarrow$  other permissions are looked at  $\Rightarrow$  perm granted

$\Rightarrow$  process can read file.

(38.)

```
int is-program(char *fname) {  
    struct stat sb;  
    if (stat(fname, &sb) == -1) return 0;  
    if (!S_ISREG(sb.st_mode)) return 0;  
    return (sb.st_mode & (S_IXUSR | S_IXGRP  
                           | S_IXOTH));
```

(39.)

```
int AplusX(const char *path) {  
    struct stat sb;  
    if (stat(path, &sb) == -1) return -1;  
    if (!S_ISREG(sb.st_mode)) return -1;  
    if ((sb.st_mode & S_IXUSR && sb.st_mode & S_IXGRP  
        && sb.st_mode & S_IXOTH) return 0;  
    if (sb.st_mode & (S_IXUSR | S_IXGRP | S_IXOTH))  
        return chmod(path, (sb.st_mode) S_IXUSR | S_IXGRP | S_IXOTH);  
    else return -1;
```

everyone  
already  
has execute  
permission

(40.) int Amicus\_w (const char \*path) {  
 struct stat sb;  
 if (stat(path, &sb) == -1) return -1;  
 no one  
 has n  
 → perm  
 to begin  
 with  
 if (! S\_ISREG(sb.st\_mode)) return -1;  
 if (!(sb.st\_mode & (S\_IWUSR | S\_IWGRP | S\_IWOTH)))  
 return 0;  
 return chmod(path, sb.st\_mode & ~ (S\_IWUSR | S\_IWGRP |  
 S\_IWOTH));

(41.) int canIread (const char \*path) {  
 struct stat sb;  
 if (stat(path, &sb) == -1) return 0;  
 uid + euid = geteuid();  
 gid + egid = getegid();  
  
 first  
 check  
 uid perm  
 if (sb.st\_uid == euid)  
 return (sb.st\_mode & S\_IRUSR);  
  
 else if (sb.st\_gid == egid)  
 return (sb.st\_mode & S\_IROGRP);  
  
 else  
 return (sb.st\_mode & S\_IROTH);

(42.)

Same as before

if define bits 8

(44.)

const int bits[3]{

int i;

unsigned char res = 0;

for (i=0; i<bits, i++) -

res = (res || (bits[i] << (bits - i - 1)))

return res;

(45.) {

while (bits->bit == 0 && bits->next)

bits = bits->next;

int res = 0;

int overflow = 0;

while (bits) {

if (overflow > 31) return -1;

res = (res << 1) | bits->bit;

b16s = b16s->next;

overflow or ft,'

3

return res;

(46.)

int i;

for (i=0; i<31; i++) {

if ((num >> (31-i)) & 1))  
    buffer[i] = '1';

else buffer[i] = '0';

}

b16sr[32] = '\0';

return buffer;

#define MAX sizeof(int) \* 8,

if (!s) return 0;

int result = 0;

int i;

int overflow;

while (s[i] == '\0') i++;

while (s[i] != '\0') {

(47.)

```

if (overflow > (MAX - 1)) return -1;
if (s[i] == '1') result = (result << 1) | 1;
else result = result << 1;
i++;
overflow++;
return result;

```

(48.)

```
count = 0;
```

```
int i;
```

```
for (i = 0; i < 32; i++) {
```

```
    if ((val >> i) & 1) count++;
}
```

```
return (count % 2);
```

```
if (count % 2)
```

(49.)

```
count = 0;
```

```
int i;
```

```
for (i = 0; i < 7; i++) {
```

```
    if ((b >> i) & 1) count++;
}
```

```
if count % 2:
```

```
b = (b) | (1 << 7);
```

$\ell(\text{se}_b = (\text{b}) \& \sim(7227),$

return b;





```
lseek(one, 5, SEEK_SET);
lseek(two,-5, SEEK_CUR);
write(one,"Twine",6);
lseek(three, 2, SEEK_END);
write(two,"glue",4);
write(three,"?",1);
```

tf: 

--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

```
lseek(two,0,SEEK_SET);
lseek(three,6,SEEK_SET);
write(one,"slip",1);
write(three,"flip?",1);
```

tf: 

--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

```
close(one);
close(two);
close(three);
```

56. () At each indicated point below, show the contents of the given file. You may assume all system calls return successfully. Even though UNIX has no end-of-file mark, you should clearly mark the current end of the file (as with  $\otimes$  below).

The initial contents of the file, "tf", are:

tf: 

W	i	n	d	o	w	s	?	$\otimes$										
---	---	---	---	---	---	---	---	-----------	--	--	--	--	--	--	--	--	--	--

```
char *filename = "tf";
int one, two, three;
```

```
one = open(filename,O_RDWR);
write(one,"For",3);
write(one,"museum",6);
```

tf: 

F	o	r	m	u	s	e	c	u	m	$\otimes$								
---	---	---	---	---	---	---	---	---	---	-----------	--	--	--	--	--	--	--	--

```
two = open(filename,O_WRONLY);
three = dup(one);
write(three,"early",2);
write(two,"Mustangs",4);
```

tf: 

M	u	s	t	u	s	e	a	m	g	$\otimes$								
---	---	---	---	---	---	---	---	---	---	-----------	--	--	--	--	--	--	--	--

```
lseek(two, 14, SEEK_SET);
lseek(three, -4, SEEK_END);
write(three,"ideal",5);
```

one ↓      two ↓

tf:	m	u	s	t	u	s	e	j	d	e	a	l	⊗						
-----	---	---	---	---	---	---	---	---	---	---	---	---	---	--	--	--	--	--	--

```
lseek(one,7,SEEK_SET);
write(three,"AREA",2);
write(two,"OS",2);
```

tf:	m	u	s	t	u	s	e	A	R	e	a	l			OS	⊗
-----	---	---	---	---	---	---	---	---	---	---	---	---	--	--	----	---

```
close(one);
close(two);
close(three);
```

```
return 0;
}
```

57. () At each indicated point below, show the contents of the given file. You may assume all system calls return successfully. Even though UNIX has no end-of-file mark, you should clearly mark the current end of the file (as with ⊗ below). **Be careful.**

The initial contents of the file, "tf", are:

tf:	T	h	i	s		o	n	e	?	⊗
-----	---	---	---	---	--	---	---	---	---	---

```
char *filename="tf";
int one, two, three;
```

```
one = open(filename,O_RDWR);
write(one,"Really?",4);
```

2/3 ↓      1 ↓

tf:	R	e	a	l	l	o	n	e	?	⊗						
-----	---	---	---	---	---	---	---	---	---	---	--	--	--	--	--	--

```
two = open(filename,O_RDWR | O_CREAT, S_IRWXU);
three = dup(two);
```

```
write(two,"It's",4);
write(three,"Traditional!",9);
```

/      2/3  
↓      ↓

tf:	I	t	'	S	T	r	a	d	i	X	i	o	n	⊗			
-----	---	---	---	---	---	---	---	---	---	---	---	---	---	---	--	--	--

```
lseek(one,-3, SEEK_END);
```

```
write(one,"Where",3);
```

```
lseek(two, 0, SEEK_SET);
```

```
write(three,"are",4);
```

```
write(two,"we?",2);
```

```
lseek(one,1, SEEK_CUR);
```

2/3  
↓

/      \  
↓      ↓

tf:	q	r	e	\	w	e	a	d	i	X	w	h	e	⊗			
-----	---	---	---	---	---	---	---	---	---	---	---	---	---	---	--	--	--

```
lseek(two,0,SEEK_SET);
```

```
lseek(three,6,SEEK_SET);
write(one,"Not ",2);
write(two,"Clever . ",3);
write(three,"Today . ",2);
```

tf: a c e l o w e l c o m e T o h e N o ☒

```
close(one);  
close(two);  
close(three);
```

58. () At each indicated point below, show the contents of the given file. You may assume all system calls return successfully. Even though UNIX has no end-of-file mark, you should clearly mark the current end of the file (as with  $\otimes$  below).

The initial contents of the file, "tf", are:

tf: M i d t e r m ?  $\otimes$

```
int one, two, three;  
char *filename="tf";
```

```
one = open(filename,O_RDWR);
write(one,"What again?",3);
```

tf: w h q t e r m ? (X)

```
two = open(filename,O_WRONLY|O_TRUNC);
```

```
three = dup(one);
```

```
write(three,"Didn't",3);
```

```
write(two,"We",3);
```

tf: We lood i d~~o~~

```
lseek(one,    13, SEEK_SET);
```

```
lseek(two,    7, SEEK_SET);
```

```
write(two,"just do",5);
```

```
lseek(three, -1, SEEK_CUR);
```

```
write(three,"this?",2);
```

```
lseek(two,5,SEEK CUR);
```

tf: w~~e~~condid just the

```
write(two,"hayfever?",2);  
write(one,"atchoo?",2);
```

tf: We did just that has @

```
close(one);  
close(two);  
close(three);
```

## networks

17.3: Declaration specifies the attributes (such as the data type) of a set of identifiers.

If the declaration also causes storage to be allocated, it is called a definition.

17.6: 1. Race between call to stat and the call to unlink, where the file can change.

2. Stat follows a symlink if symbolic points to a regular file we would remove the symlink and not the socket file. → use lstat.

## Signals:

10.7: pause returns -1 and program exits (errno is set to EINTR)  
→ pause returns whenever a signal is caught.

10.3:

## Process Lifecycle:

Chapter 7:

7.2:

## Process fork() & wait()

8.5:

MYARG2

8.6:

main

:

if (fork() == 0);

exit(0);

exit(0);

(64)

```
int make_new_log (char *fname) {
```

struct stat sb;

```
if (stat (fname, &sb) > -1) return -1;
```

```
return open (fname, O_WRONLY, (S_IWRITE | S_INO)) ;
```

## Unreliable Signal Handling:

→ signal handler is called with the signal number

→ (maybe) the handler is reset to SIG\_DFL

→ (maybe) the further deliver of the same signal is blocked

→ (maybe) soon system calls are interrupted, returning error error and setting errno to EINTR

→ reliable signal handling. All this is done with certainty

→ provides support for more timing, but makes the interface more complicated.

Task: set up SIGINT & blocking over a critical section.

```
void sigint_handler(int num){  
    printf("Hello :)\n");
```

3

```
int main(int argc, char * argv[]){
```

```
    sigset(SIGINT, new, old);
```

```
    struct sigaction sa;
```

```
    sa.sa_flags = 0;
```

```
    sigemptyset(&sa.sa_mask);
```

```
    sa.sa_handler = sigint_handler;
```

```
    sigemptyset(&new);
```

```
    sigaddset(&new, SIGINT);
```

```
    sigprocmask(SIG_BLOCK, &new, &old);
```

```
    if(sigaction(SIGINT, &sa, NULL) == -1)
```

```
        perror("sigaction");  
        exit(1);
```

```
    printf("Important code !\n");
```

```
    sigprocmask(SIG_SETMASK, &old, NULL);
```

```
    return 0;
```

3

# Rockit

```
int main (int argc, char *argv[]) {
    printf ("Hello world!\n");
    pid_t child = fork();
    if (child == -1) {
        perror ("fork");
        exit (EXIT_FAILURE);
    }
    else if (child == 0) {
        printf ("PID: %d\n", getpid());
        exit (EXIT_SUCCESS);
    }
    else {
        printf ("PID: %d\n", getpid());
        if (wait (WIFCL) == -1) {
            perror ("wait");
            exit (EXIT_FAILURE);
        }
        printf ("Bye!\n");
    }
    return 0;
}
```

```
int main( int argc, char * argv ) {  
    /* command line */  
    char * fname;  
    pid_t child = fork();  
    int status;  
    if ( child == -1 ) {  
        perror("fork");  
        exit(EXIT_FAILURE);  
    } else if ( child == 0 ) {  
        exec( fname, fname, (char *) NULL );  
        perror("exec");  
        exit(EXIT_FAILURE);  
    }  
}
```

```
else {  
    if ( wait(&status) == -1 ) {  
        perror("wait");  
        exit(EXIT_FAILURE);  
    }  
}
```

```
if ( WIFEXITED(status) == 0 ) {
```

```
    printf("child died with status %d\n",  
           WEXITSTATUS(status));  
}
```

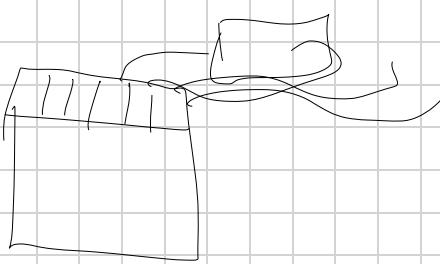
```
else {  
    if ( wait(&status) == -1 ) {  
        perror("wait");  
    }  
}
```

```
else {
```

points (abnormally)

```
int p-one[2];
```

```
int p-two[2];
```



```
pid_t child1, child2, child3;
```

```
; if (child1 == fork()) == -1 {
```

```
    perror("fork");
```

```
    exit(1);
```

```
else if (child1 == 0) {
```

```
; if (dup2(p-one[1], STDOUT_FILENO) == -1) {
```

```
    perror("dup2");
```

```
    exit(1);
```

```
close(p-one[1]);
```

```
close(p-one[0]);
```

```
close(p-two[1]);
```

```
close(p-two[0]);
```

```
execp("ls") ("ls", (char *)NULL);
```

```
perror("ls");
```

```
exit(1);
```

}

if ( $child(d2) = false()$ ) = -1)

    Penosfah  
    Exit(1);

}

else if ( $child(d2) = 0$ )

    if ( $dup2(p-one[0], STDIN_FILENO) = -1$ ) {  
        Error

}

    if ( $dup2(p-two[1], STDOUT_FILENO) = -1$ ) {  
        Error

}

    close(p-one[1]);

    close(p-one[0]);

    close(p-two[1]);

    close(p-two[0]);

    execvp("sort", {"sort", "-r", (char \*)NULL});  
    perror("sort");  
    exit(1);

if ( $child(d3) = false()$ ) = -1) { error

else ; if ( $child(d) = 0$ ) {

    if ( $dup2(p-two[0], STDIN_FILENO) = -1$ ) {

        Error

close (p-one[1]);

close (p-one[0]);

close (p-two[1]);

close (p-two[0]);

execp ("move", "more", (char \*)NULL);  
perror ("more");  
exit(1);

}

(\* parent \*)

close (p-one[1]);

close (p-one[0]);

close (p-two[1]);

close (p-two[0]);

for (int i=0; i<3; i++)

if (wait (NULL) == -1)

perror

{}

exit(0);

struct pollfd

int fd;

int events;

int revents;

3

int poll(struct pollfd \*fd, int length time  
timeout)

struct pollfd fd[2];

fd[0].fd = STDIN\_FILENO

fd[0].events = POLLIN;

fd[0].revents = 0;

fd[1] = fd[0];

fd[1].fd = network\_socket;

for(;;) {

poll(fd, 2, -1);

if (fd[0].revents & POLLIN) {

/\* stdin \*/

3 if (fd[1].revents & POLLIN) {

/\* handle network \*/

3

## Networking :

### TCP :

#### Server :

```
int sockfd = socket(AF_INET, SOCK_STREAM, 0);
```

```
struct sockaddr_in sa, newsockinfo, peer;
```

```
sa.sin_family = AF_INET;
```

```
sa.sin_port = htons(port);
```

```
sa.sin_addr.s_addr = htonl(INADDR_ANY);
```

```
bind(sockfd, (struct sockaddr*)&sa, sizeof(sa));
```

```
int newsock = accept(sockfd, (struct sockaddr*)&peer, &sizeof(peer));
```

65.

```
#define LEN 706
```

```
char * allocPath();
```

```
char * path = NULL;
```

```
int len = SIZE;
```

```
while (1) {
```

```
path = (char *) realloc(path, len + sizeof(char));
```

```
path = getcmd(path, len);
```

```
if (path)
```

```
break;
```

```
else if (errno != ERANGE && errno != 0)
```

```
return NULL;
```

```
}
```

```
len += SIZE
```

```
}
```

```
path = realloc(path, strlen(path) + 1);
```

```
return path
```

(66) int isroot(const char \*path) {

```
struct stat sroot, spath;
```

```
if (stat("/", &sroot) == -1) return 0;
```

```
if (stat(path, &spath) == -1) return 0;
```

return ((s.root.st\_dev == s.path.st\_dev) &&  
(s.root.st\_ino == s.path.st\_ino)),

(67.) int is\_leaf\_directory (const char \*path) {

struct stat sb; int res = 1;

if (stat(path, &sb) == -1) return 0;

if (!S\_ISDIR(sb.st\_mode)) return 0;

char cwd [PATH\_MAX + 1];

getcwd(cwd, PATH\_MAX);

if (!cwd) return 0;

cwd [PATH\_MAX] = '\0';

if (chdir(path) == -1) return 0;

DIR \*current;

struct dirent \*dirb;

current = opendir("."),

if (!current) return 0;

while ((dirb = readdir(current)) != NULL) {

if (!strcmp(dirb.d\_name, ".") &&

!strcmp(dirb.d\_name, "..")) continue;

if (stat(dirb.d\_name, &sb) == -1) res = 0;

if (S\_ISDIR(sb.st\_mode)) res = 0;

```
if (closedir(current) == -1) res = 0;  
if (chdir(curl) == -1) res = 0;
```

```
return res;
```

3

(68.)

```
int rm_tree(const char *victim){  
    struct stat sb;  
    if (stat(victim, &sb) == -1) return 0;  
    if (!S_ISDIR(sb.st_mode)) {  
        if (unlink(victim) == -1) return 0;  
        return 1;  
    }
```

3

```
else {
```

```
    char cwd[PATH_MAX + 7];
```

```
    getcwd(cwd, PATH_MAX);
```

```
    if (!curl) return 0;
```

```
    cwd[PATH_MAX] = '\0';
```

```
    if (chdir(victim) == -1) return 0;
```

```
    DIR *current;
```

```
    if (opendir("..") == -1) return 0;
```

```
    int res = 1;
```

```
    struct dirent *dир;
```

```
    while ((дир = readdir(current)) != NULL) {
```

if (!strcmp (dirp->d\_name, "") ||  
!strcmp (dirp->d\_name, "..")) continue;

if (stat (dirp->d\_name, &sb) == -1) res = 0;

if (S\_ISDIR (sb.st\_mode)) {

if (!rm\_tree (dirp->d\_name)) res = 0;

} else {

if (unlink (dirp->d\_name) == -1) res = 0;

}

if (closedir (current) == -1) res = 0;

if (closedir (curr) == -1) res = 0;

if (rmdir (victim) == -1) res = 0;

return res;

}

~

(68) int countfiles (const char \*path) {

```

int result = 0;
struct stat sb;
if (lstat (path, &sb) == -1) return -1;
if (S_ISREG (sb.st_mode)) return 1;
if (S_ISDIR (sb.st_mode)) {
    int buf = 0;
    char cwd [PATH_MAX];
    getcwd (cwd, PATH_MAX);
    if (! cwd) return -1;
    if (chdir (path) == -1) return -1;
    DIR *current;
    current = opendir (".");
    if (! current) return -1;
    struct dirent *dirp;
    while ((dirp = readdir (current)) != NULL) {
        if (! strcmp (dirp->d_name, ".") ||
            ! strcmp (dirp->d_name, "..")) continue;
        if (lstat (dirp->d_name, &sb) == -1) return -1;
        if (S_ISREG (sb.st_mode)) result++;
        if (S_ISDIR (sb.st_mode)) {
            if ((buf = countfiles (dirp->d_name)) == -1) return -1;
            result += buf;
        }
    }
}

```

}

if (closedir (current) == -1) return -1;  
if (chdir (cud) == -1) return -1;  
return result;

3

return 0;

}

## Signal handling:

```
void sigintHandler(int num) {
```

```
    printf("Neener neener \\n^4"),
```

```
}
```

```
int main(int argc, char *argv[]) {
```

```
    struct sigaction sa, sold;
```

```
    sa.sa_handler = sigintHandler;
```

```
    if (sigemptyset(&sa.sa_mask) == -1)
```

```
        perror("sigemptyset");
```

```
        exit(EXIT_FAILURE);
```

```
}
```

```
    sa.sa_mask = 0;
```

```
    if (sigaction(SIGINT, &sa, &sold) == -1)
```

```
        perror("sigaction");
```

```
        exit(EXIT_FAILURE);
```

```
}
```

```
    if (sigset(SIGINT, newmask, oldmask))
```

```
        if (sigemptyset(&newmask) == -1) { perror("sigemptyset")
```

```
            exit(1);
```

```
}
```

```
        if (sigaddset(&newmask, SIGINT) == -1)
```

```
perror("sigaddset"),  
exit(EXIT_FAILURE);
```

{

```
if (sigprocmask(SIG_BLOCK, &newmask, &oldmask)  
== -1)  
perror("sigprocmask"),  
exit(EXIT_FAILURE);
```

}

```
printf("Important stuff!\\n");
```

```
if (sigprocmask(SIG_SETMASK, &oldmask, NULL)  
== -1){  
perror("sigprocmask"),  
exit(EXIT_FAILURE);
```

}

```
return 0;
```

}