

10 Lecture: Unbuffered IO, cont.

Outline:

- Announcements
- Unix Overview
- Identity Issues: logging in
 - Looking at system files
- From Last, Last, Last Time: Unix Overview
- Files and Directories
 - Directories
 - Directory Manipulation
- System Calls
- From last time: Files and the filesystem
- Basic File IO
 - open(2)
 - creat(2)
 - close(2)
 - read(2)
 - write(2)
- Performance: Buffered vs. Unbuffered
- Review: Unbuffered IO
- Onwards: lseek(2)
- Next Time
- If there's time: Lab03/Asgn3
 - The assignment
- From Email: Huffman
 - Huffman Codes
 - Reminder: Setting and clearing bits

10.1 Announcements

- Coming attractions:

Event	Subject	Due Date			Notes
lab03	htable	Fri	Oct 27	23:59	
asgn3	hencode/hdecode	Fri	Nov 3	23:59	
lab05	mypwd	Mon	Nov 6	23:59	
asgn4	mytar	Mon	Nov 27	23:59	
asgn5	mytalk	Fri	Dec 1	23:59	
lab07	forkit	Mon	Dec 4	23:59	
asgn6	shell	Fri	Dec 8	23:59	

Use your own discretion with respect to timing/due dates.

- Back from the dead. Catching up:
 - We have 18 classes left. We can do this. I counted; It fits.
 - Weights in Gradesheet Snapshot are currently meaningless
 - No asgn2 (I'll post it if you want to play)

- No C Quiz (old questions are on the web if you want to challenge yourself)
- Do we need to talk more about Make or are you good?
- Lab02 and asgn1 will close Monday. It's time to move on
- Other dates will happen soon (famous last words)
- The “Gradesheet snapshot”.
- gprof(1)
- Assignments: Hours spent are irrelevant.
 - Effective hours
 - Decomposition
 - Incremental development
 - Incremental testing
 - avoid the Death March mentality
- I gave you test cases. It seems almost nobody ran them. This would have solved UI problems.
- qsort() demo
- there's stuff in the notes that's not always in class

10.2 Unix Overview

When discussing an operating system, everything depends on everything else. Today we want to talk about a process's view of its world in terms of various services an OS provides.

A note on the examples: They use code contained in Appendix B.

<http://www.kohala.com/start/apue.html>

A process's view of the world:

- Identity (uid, gid)
- Filesystem (storage, organization, ownership, access)
- IO (What good is a filesystem if you can't use it)
- Programs and Processes
- Interprocess Communication: Signals

10.3 Identity Issues: logging in

Before doing anything else on a unix machine, you need to work out who you are. (User and group identities)

Identity consists of uid and gid

Identity is controlled by several system files:

- `/etc/passwd` — this maps login name to uid (and stores passwd)
- `/etc/group` — primary and supplemental groups

- shadow passwd: `/etc/shadow`
- yellowpages (or other directory schemes, LDAP, etc.): shared info across machines

A password line:

```
csc357:*:1659:350:csc357:/home/cscstd/abcd/csc357:/bin/tcsh
```

This is “login:passwd:uid:gid:comment:home dir:shell”

How the password algorithm works: Your typed password and the salt (the first two characters) are combined via a cryptographic hash to produce 13 printable characters in “[a-zA-Z0-9./]”.

```
#include <unistd.h>
```

```
char *crypt(const char *key, const char *salt);
```

Traditionally a variation on DES, now likely to be something else..

The salt perturbs the dictionary 4096 different ways.

Your uid is your identity. The password file is the only place your login name appears in the system.

This exposure makes them vulnerable.

10.3.1 Looking at system files

real vs. shadow or yp

shadow/yp on hornet and drseuss

10.4 Files and Directories

The entire system is connected through the filesystem. (discussion of limits.h)

NAME_MAX:

Solaris: 512

linux: 256

10.4.1 Directories

A directory is a file containing a list of:

- name
- attributes (not actually in the directory)
- where (inode)

described in “dirent.h”:

```
struct dirent
{
    long d_ino;           /* inode number */
    off_t d_off;          /* offset to this dirent */
    unsigned short d_reclen; /* length of this d_name */
    char d_name [NAME_MAX+1]; /* file name (null-terminated) */
}
```

Directory entries are unordered.

Directories are protected.

10.4.2 Directory Manipulation

Manipulation functions `opendir(3)`, `readdir(3)`, `closedir(3)`, `rewinddir(3)`

```
#include <sys/types.h>
#include <dirent.h>
DIR *opendir(const char *name);
int closedir(DIR *dir);
struct dirent *readdir(DIR *dirp);
void rewinddir(DIR *dir);
```

As part of its environment, every process has a:

Home Directory defined in `/etc/passwd`

Current Working Directory starts at home and follows `chdir()`

10.5 System Calls

Before we talk about IO services...

Look like C functions, but are direct requests for OS services.

A system call is an entry into the kernel.

Linux: (RH7.0) 222
Solaris: 253
Minix: 53

10.6 From last time: Files and the filesystem

- A file is a collection of stuff
- Files go into the filesystem, managed by Directories

10.7 Basic File IO

The basic file IO functions can be handled by the following system calls:

```
int open(const char *pathname, int flags, mode_t mode);
ssize_t read(int fd, void *buf, size_t count);
ssize_t write(int fd, const void *buf, size_t count);
int close(int fd);
off_t lseek(int fildes, off_t offset, int whence);
```

10.7.1 `open(2)`

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
```

```
int open(const char *pathname, int flags);
int open(const char *pathname, int flags, mode_t mode);
int creat(const char *pathname, mode_t mode);
```

`flags` is one of `O_RDONLY`, `O_WRONLY` or `O_RDWR` which request opening the file read-only, write-only or read/write, respectively.

`O_CREAT`

If the file does not exist it will be created.

`O_TRUNC`

If the file already exists it will be truncated.

`O_APPEND`

The file is opened in append mode. [Before each write the pointer is moved to the end]

Symbolic names are supplied. 0600, e.g. is `S_IRUSR | S_IWUSR`

S_IRWXU	<code>S_IWGRP</code>
<code>S_IRUSR</code>	<code>S_IXGRP</code>
<code>S_IWUSR</code>	S_IRWXO
<code>S_IXUSR</code>	<code>S_IROTH</code>
S_IRWXG	<code>S_IWOTH</code>
<code>S_IRGRP</code>	<code>S_IXOTH</code>

Man `stat(2)` on linux; `stat(3HEAD)` on solaris. (Go figure)

`open(2)` returns a valid file descriptor or `-1` and sets `errno`.

<code>ENOENT</code>	given file name doesn exist
<code>ENOTDIR</code>	a component is not a directory
<code>EFAULT</code>	bad address passed for name
<code>EACCESS</code>	Permission denied
<code>ENOMEM</code>	out of memory
<code>ENAMETOOLONG</code>	file name is too long

10.7.2 `creat(2)`

SYNOPSIS

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>
```

```
int creat(const char *pathname, mode_t mode);
```

Equivalent to `open(path, O_WRONLY | O_CREAT | O_TRUNC, mode)`

10.7.3 `close(2)`

```
#include <unistd.h>
```

```
int close(int fd);
```

RETURN VALUE

`close()` returns zero on success. On error, `-1` is returned, and `errno` is set appropriately.

ERRORS

`EBADF` `fd` isn't a valid open file descriptor.

`EINTR` The `close()` call was interrupted by a signal; see `signal(7)`.

`EIO` An I/O error occurred.

10.7.4 read(2)

SYNOPSIS

```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buf, size_t count);
```

(`ssize_t` is signed. `size_t` is unsigned)

Returns number of bytes read or `-1` on error..

10.7.5 write(2)

SYNOPSIS

```
#include <unistd.h>
```

```
ssize_t write(int fd, const void *buf, size_t count);
```

Returns number of bytes written or `-1` on error..

10.8 Performance: Buffered vs. Unbuffered

The issue of whether to use buffered IO or not often depends on performance.

```
% ll KeepArchive.tar.gz
-rw----- 1 pnico pnico 51163086 Mar 20 2001 KeepArchive.tar.gz
```

```
% /usr/bin/time mycp KeepArchive.tar.gz outfile
37.29user 7.09system 0:50.99elapsed 87%CPU
```

```
% /usr/bin/time unbuffered KeepArchive.tar.gz outfile 1
47.64user 260.28system 5:12.10elapsed 98%CPU
```

```
% /usr/bin/time unbuffered KeepArchive.tar.gz outfile 4096
0.32user 19.27system 0:29.83elapsed 65%CPU
```

```
% /usr/bin/time unbuffered KeepArchive.tar.gz outfile 8192
0.23user 21.72system 0:28.94elapsed 75%CPU
```

```
% /usr/bin/time unbuffered KeepArchive.tar.gz outfile 65536
0.10user 23.70system 0:29.08elapsed 81%CPU
```

```
% /usr/bin/time unbuffered KeepArchive.tar.gz outfile 131072
0.04user 23.97system 0:28.80elapsed 83%CPU
```

For comparison's sake:

```
% /usr/bin/time /bin/cp KeepArchive.tar.gz outfile
0.58user 20.54system 0:30.86elapsed 68%CPU
```

```

#include <stdio.h>

void stdio_cp(char *iname, char*outname){
    FILE *infile,*outfile;
    int c;

    if ( NULL == (infile = fopen(iname,"r")) ) {
        perror(iname);
        exit(-1);
    }
    if ( NULL == (outfile = fopen(outname,"w")) ){
        perror(outname);
        exit(-1);
    }

    while((c=getc(infile))!=EOF) {
        if ( putc(c,outfile) == EOF ) {
            perror("putc");          /* putc returns EOF on error */
            exit(-1);
        };
    }

    fclose(infile);
    fclose(outfile);
}

```

Figure 51: cp based on stdio


```

#include <unistd.h>
#include<stdlib.h>
#include<string.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

void unbuffered_cp(char *iname, char* outname, int size){
    int infd, outfd;
    int num;
    char *buffer;

    buffer = (char*)safe_malloc(size);

    if ( (infd = open(iname, O_RDONLY)) < 0 ) {
        perror(iname);
        exit(-1);
    }
    if ( (outfd = open(outname,
        (O_WRONLY | O_CREAT | O_TRUNC),
        (S_IRUSR | S_IWUSR))) < 0 ) {
        perror(outname);
        exit(-1);
    }

    while((num=read(infd,buffer,size)) > 0){
        if ( write(outfd,buffer,num) != num){
            perror("write");
            exit(-1);
        }
    }
    if ( num < 0 ) {
        perror("read");
        exit(-1);
    }

    if ( close(infd) ) {
        perror("close");
        exit(-1);
    }
    if ( close(outfd) ) {
        perror("close");
        exit(-1);
    }

    free(buffer);
}

```

Figure 52: cp based on unbuffered IO

10.9 Review: Unbuffered IO

```
int open(const char *pathname, int flags, mode_t mode);
ssize_t read(int fd, void *buf, size_t count);
ssize_t write(int fd, const void *buf, size_t count);
int close(int fd);
off_t lseek(int fildes, off_t offset, int whence);
```

10.10 Onwards: lseek(2)

Move the file pointer (or find out where it is). This movement has no effect on the file until the next write.

`lseek(2)` introduces the concept of “holes”, places in the file where nothing has ever been written. These locations are read as `'\0'`.

SYNOPSIS

```
#include <sys/types.h>
#include <unistd.h>

off_t lseek(int fildes, off_t offset, int whence);
```

DESCRIPTION

The `lseek` function repositions the offset of the file descriptor `fildes` to the argument `offset` according to the directive `whence` as follows:

SEEK_SET

The offset is set to `offset` bytes.

SEEK_CUR

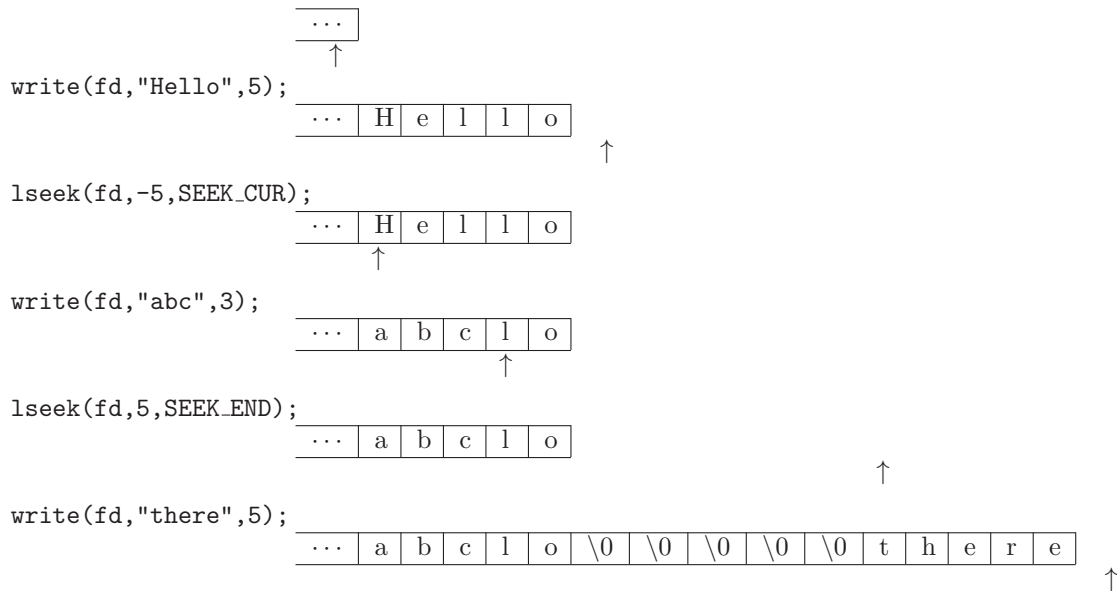
The offset is set to its current location plus `offset` bytes.

SEEK_END

The offset is set to the size of the file plus `offset` bytes.

Examples:

<code>lseek(fd, 0, SEEK_SET)</code>	rewind the file
<code>lseek(fd, 0, SEEK_END)</code>	get ready to append
<code>pos = lseek(fd, 0, SEEK_CUR)</code>	get current location



10.11 Next Time

C language quiz:

Quiz next time: K&R 1–7: (Yes, actually read it.)

- Types, operators, expressions
- Control flow
- Functions (and macros)
- Pointers and Arrays
- Complex data: structures, typedef, etc.
- Memory Management
- IO (stdio)
- **not** chapter 8

My usual exam style is short-answer questions for breadth and a couple of longer ones to demonstrate depth of knowledge. I try to emphasize understanding over rote knowledge, but bear in mind that in a field such as this two are hard to separate.

- Short answer for breadth, longer answer for depth
- Don't worry about cramming the whole C library into your head. Useful things will show up at the end of the exam and/or be described.
- There are sample C quizzes linked from the web page

10.12 If there's time: Lab03/Asgn3

10.12.1 The assignment

Discussion of the assignment:

Write two programs to compress and uncompress files:

Usage:

hencode infile [outfile]

hdecode [(infile | -) [outfile]]

10.13 From Email: Huffman

Right, you're only padding the last byte with bits. My concern about endianness is this: If you shift bits into an int, the bytes will wind up in the wrong order on a little-endian machine. Consider:

```
int x = 0;
x |= 0x0A << 24;
x |= 0x0B << 16;
x |= 0x0C << 8;
x |= 0x0D << 0;
printf("0x%08x\n", x);
```

This will print 0x0A0B0C0D because that's the integer, but, if you write this to the disk using "write(fd, &x, sizeof(int);", the bytes on the disk will be 0x0D 0x0C 0x0B 0x0A. This is because on a little-endian machine the least significant end goes first. This is why I say build it a byte at a time so there's no confusion about endianness. You can build it into an array of chars and write the whole array (or at least a hunk of it) at a time, but build it as bytes, not integers.

10.13.1 Huffman Codes


David A. Huffman. A method for the construction of minimum-redundancy codes. *Proceedings of the Institute of Radio Engineers*, 40(9):1098–1101, September 1952.

Building the tree


```
% cat test
aabbccddd
%
```

Byte	Count
0x0a	'\n' 1
0x61	'a' 2
0x62	'b' 2
0x63	'c' 2
0x64	'd' 3


1. Building the tree: Stage 0




2. Building the tree: Stage 1




3. Building the tree: Stage 2



4. Building the tree: Stage 3



5. Building the tree: Stage 4



File Format

Num - 1	c1	count of c1	c2	count of c2	...	cn	count of cn
(uint8_t)	(uint8_t)	(uint32_t)	(uint8_t)	(uint32_t)	...	(uint8_t)	(uint32_t)
1 byte	1 byte	4 bytes	1 byte	4 bytes	...	1 byte	4 bytes

...	... encoded characters
-----	----------------------------	-----

Encoding the File

Byte	Code
0x0a '\n'	100
0x61 'a'	101
0x62 'b'	00
0x63 'c'	01
0x64 'd'	11

File: "aabbccddd\n"

0000 0100	0000 1010	0000 0000	0000 0000	0000y0000	0000 0001	011 00001	0000 0000
04	0a	00	00	00	01	61	00
0000 0000	0000 0000	00000010	01100010	0000 0000	0000 0000	0000 0000	00000010
00	00	02	62	00	00	00	02
0110 0011	0000 0000	0000 0000	0000 0000	0000 0010	0110 0100	0000 0000	0000 0000
63	00	00	00	02	64	00	00
0000 0000	0000 0011						
00	03						
1011 0100	0001 0111	1111 1000					
b4	17	f8					

Results

```
% hencode test test.huff
% ls -l test test.huff
-rw----- 1 pnico pnico 10 Oct 17 09:55 test
-rw----- 1 pnico pnico 32 Oct 17 12:19 test.huff
% od -tx1 test.huff
00000000 04 0a 00 00 00 01 61 00 00 00 02 62 00 00 00 02
00000020 63 00 00 00 02 64 00 00 00 03 b4 17 f8
00000035
```

10.13.2 Reminder: Setting and clearing bits

It's important to remember how to set and clear bits in a bitfield while preserving the rest of the bits:

Given a word, *word*, and a bitmask *mask*

- To Set Bits:

Bitwise-OR the word with the mask:

$$word \vee mask$$

Bitwise, anything ORed with 1 is set and ORed with 0 is itself:

	<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>
OR	1	0	1	1	0
<hr/>					
	1	<i>b</i>	1	1	<i>e</i>

- To Clear Bits:

Bitwise-AND the word with the inverse of mask:

$$word \wedge \overline{mask}$$

Bitwise, anything ANDed with 1 is itself and ANDed with 0 is 0:

	<i>a</i>	<i>b</i>	<i>c</i>	<i>d</i>	<i>e</i>
AND	1	0	1	1	0
<hr/>					
	<i>a</i>	0	<i>c</i>	<i>d</i>	0

- In C:

Bitwise OR: |
Bitwise AND: &
Bitwise NOT: ~
Bitwise XOR: ^

An example in Figure 53.

```
void print_bits(unsigned char c) {  
    unsigned char mask;  
    for(mask=0x80; mask; mask>>=1)  
        if ( mask & c )  
            putchar('1');  
        else  
            putchar('0');  
}
```

Figure 53: printing the bits in a byte.